


An Operational Method Toward Efficient Walk Control Policies for Humanoid Robots

Ludovic Hofer, Quentin Rouxel

ludovic.hofer@labri.fr, quentin.rouxel@labri.fr

Problem Definition

Humanoid robots playing soccer



Problem: Approach the ball and kick:

- Drive the walk to place the robot in kick position
- Align with goal posts orientation
- Do not touch the ball
- On holonomic robot (HA)
- On almost non holonomic robot (ANHA)

Action space (dimension = 3): Walk is controlled in acceleration.

Name	Units	min	max
Forward acc	$\frac{m}{step^2}$	-0.02	0.02
Lateral acc	$\frac{m}{step^2}$	-0.01	0.01
Angular acc	$\frac{rad}{step^2}$	-0.15	0.15

State space (dimension = 6):

Name	Units	min	max
Ball distance	m	0	1
Ball direction	rad	$-\pi$	π
Kick direction	rad	$-\pi$	π
Forward speed	$\frac{m}{step}$	-0.02	0.04
Lateral speed	$\frac{m}{step}$	-0.02	0.02
Angular speed	$\frac{rad}{step}$	-0.2	0.2

Problem features:

- Continuous action space
- Continuous state space
- Stochastic displacement model
- Discontinuous reward

Compared Policies

Winner2016: Expert policy used by Rhoban team to win RoboCup 2016. Manually tuned parameters.

CMA-ES: Same as Winner2016 with parameters tuned using black-box optimization CMA-ES in simulation.

RFPI: Policy represented as regression forest. Computed using MDP Random Forest Policy Iteration (RTFI) solver.

Proposed Method

1. Learn displacement model from data
2. Solve MDP problem in simulation
3. Apply on real robot

Displacement Model Learning

- Large discrepancies between walk orders and actual displacement
- Use data to learn corrective linear model
- No need for external hardware


$$\Delta(x, y, \theta)_{\text{walk orders}, k} \mapsto \Delta(x, y, \theta)_{\text{corrected}, k}$$

Compare different linear models:


$$\begin{bmatrix} \Delta x_{\text{corrected}} \\ \Delta y_{\text{corrected}} \\ \Delta \theta_{\text{corrected}} \end{bmatrix} = M \begin{bmatrix} 1 \\ \Delta x_k \\ \Delta y_k \\ \Delta \theta_k \end{bmatrix}$$

- Proportional Model: 3 parameters.
- Simple Linear Model: 6 parameters.
- Full Linear Model: 12 parameters.

Parameter learning through black-box (CMA-ES) optimization with 20 learning data sequences:



Model parameters convergence: (full linear model):



Contributions

- MDP continuous state and action solver (random forests)
- Displacement model learning procedure without external hardware
- Real robot applications
- Approach time improved on holonomic and non holonomic humanoid robot

MDP RFPI Solver

The CSA-MDP learner algorithm:

```

1:  $\pi = \text{getRandomPolicy}()$ 
2:  $V = \text{buildConstantApproximator}(0)$ 
3:  $\text{visitedStates} = \text{seedStates}$ 
4:  $\text{policyId} = 1$ 
5:  $\text{runId} = 0$ 
6: while  $\text{timeRemaining}()$  do
7:    $\text{executeRun}(\pi, \text{visitedStates})$ 
8:    $\text{runId}++$ 
9:   if  $\text{runId} == \text{policyId}$  then
10:    // Perform roll outs from visited states and
11:    // fit a regression forest with piecewise
12:    // constant model approximators.
13:     $V = \text{updateValue}(\pi, V, \text{visitedStates})$ 
14:    // For every visited state: 1-step
15:    // optimization of action.
16:    // Then fit a regression forest with
17:    // piecewise linear model approximators.
18:     $\pi = \text{updatePolicy}(V, \text{visitedStates})$ 
19:     $\text{runId} = 0$ 
20:     $\text{policyId}++$ 
21:   end if
22: end while

```

Results


In simulation: average fitness costs:

	Winner2016	CMA-ES	RFPI
HA	31.84	14.90	11.88
ANHA	44.12	36.18	15.97

On physical robot: average time in seconds before kicking the ball:

	Winner2016	CMA-ES	RFPI
HA	19.98	13.72	11.45
ANHA	48.14	25.69	18.81

Typical trajectories:



References

- Hofer, L., and Gimbert, H. Online reinforcement learning for real-time exploration in continuous state and action markov decision processes. In PlanRob2016, ICAPS 2016.
- Rouxel, Q.; Passault, G.; Hofer, L.; N'Guyen, S.; and Ly, O. Learning the odometry on a small humanoid robot. In ICRA 2016.
- Ernst, D.; Geurts, P.; and Wehenkel, L. Tree-Based Batch Mode Reinforcement Learning. In JMLR 2005.